Table 8. Prediction accuracy by Baseline tier and condition.

Baseline tier	Control (%)	Incentive (%)	Difference (pp)
Current-generation	91.6	91.2	−0.4
Prior-generation	66.6	71.8	+5.2

Note. Current-generation = Gemini 2.5 Pro, Claude Sonnet 4.5, DeepSeek R2. Prior-generation = Gemma 2 2B IT, Claude 3 Haiku, DeepSeek R1. pp = percentage points.

- **Current-generation Baselines** (Gemini 2.5 Pro, Claude Sonnet 4.5, DeepSeek R2)
- **Prior-generation Baselines** (Gemma 2 2B, Claude 3 Haiku, DeepSeek R1)

What's the main thing driving that? Overall, the Baseline models were really good at getting the questions correct. If you're always going to get the questions correct, it's a safe bet to assume you'll get it right.

Table 9. Baseline model accuracy on math/logic questions.

Model	Generation	Overall Correct (%)	Yes/No Questions (%)	Numeric Questions (%)
1.Google Gemini 2.5 Pro	Current	95.0	88.2	98.5
2.Claude Sonnet 4.5	Current	90.0	85.3	92.4
3.DeepSeek R2	Current	95.0	88.2	98.5
4.Google Gemma 2 2B IT	Prior	22.0	47.1	9.1
5.Claude 3 Haiku	Prior	38.0	44.1	34.8
6.DeepSeek R1	Prior	97.0	91.2	100.0

Note. Current-generation models accessed via private web interfaces; prior-generation via web interfaces or HuggingFace Chat.

Note that DeepSeek R1, while preregistered as a prior-generation model, achieved 97% accuracy and behaved more like the current-generation models. To maintain our consistency with pre-registration (where we classified R1 as prior-generation) we left it as prior generation in all analysis. This makes the prior tier heterogeneous and complicates any strong claims about *generation* per se; the main signal is really about forecasting high- vs low-performing baselines.

A statistical pattern quickly emerged. If the “smart” models were almost always getting the questions right, the smart money is to bet on them getting the answers right. This is what we saw after in both the control and incentive models. The table below shows how often the Predictor models would assume a Baseline model would get the question correct. And many of the cases assumed it would be 100% of the time.

Table 10. Percentage of predictions where Predictor assumed Baseline would be correct.

Run	Cond.	Gemini	Sonnet	R2	Gemma	Haiku	R1	Avg
Grok-1	Cont	96.0	94.0	98.0	46.0	46.0	98.0	79.7
Grok-2	Cont	96.0	94.0	94.0	48.0	53.0	95.0	80.0
Grok-3	Cont	96.0	92.0	99.0	51.0	56.0	98.0	82.0
Grok-1	Ince	100.0	100.0	100.0	39.0	39.0	100.0	79.7
Grok-2	Ince	100.0	100.0	100.0	37.0	37.0	100.0	79.0
Grok-3	Ince	100.0	100.0	100.0	35.0	38.0	100.0	78.8
Mistral-1	Cont	100.0	100.0	100.0	52.0	52.0	37.0	73.5
Mistral-2	Cont	100.0	100.0	100.0	26.0	26.0	100.0	75.3
Mistral-3	Cont	100.0	100.0	100.0	51.0	51.0	100.0	83.7
Mistral-1	Ince	100.0	81.0	78.0	46.0	46.0	91.0	73.7
Mistral-2	Ince	100.0	100.0	100.0	20.0	20.0	100.0	73.3
Mistral-3	Ince	75.0	100.0	100.0	0.0	0.0	100.0	62.5
Qwen-1	Cont	87.0	86.0	90.0	53.0	52.0	89.0	76.2
Qwen-2	Cont	98.0	96.0	93.0	51.0	53.0	97.0	81.3
Qwen-3	Cont	96.0	83.0	95.0	50.0	41.0	97.0	77.0
Qwen-1	Ince	100.0	100.0	100.0	15.0	15.0	100.0	71.7
Qwen-2	Ince	100.0	100.0	100.0	50.0	75.0	100.0	87.5
Qwen-3	Ince	100.0	100.0	100.0	20.0	20.0	91.0	71.8
	Control Total	**96.6**	**93.9**	**96.6**	**47.6**	**47.8**	**90.1**	**78.8**
	Incentive Total	**97.2**	**97.9**	**97.6**	**29.1**	**32.2**	**98.0**	**75.3**

Note. Values show "Yes" rate—the percentage of the 100 questions where the Predictor predicted the Baseline would answer correctly. Gemini = Gemini 2.5 Pro; Sonnet = Claude Sonnet 4.5; R2 = DeepSeek R2; Gemma = Gemma 2 2B IT; Haiku = Claude 3 Haiku; R1 = DeepSeek R1. Cond. = Condition (Cont = Control, Ince = Incentive). Avg = average across all 6 Baselines.

Two qualitative patterns emerge:

1. **Large tier gap**
   Predictors are *extremely* good at forecasting high-performing current-generation models, which are themselves near the ceiling on this dataset. For these Baselines, it is relatively easy for a Predictor to identify and bet heavily on "sure things." This pattern is consistent with H3, but given the small sample and descriptive analysis, we treat it as exploratory.

2. **Asymmetry on hard items**
   For structurally harder items, Predictors are mediocre at forecasting whether current-generation Baselines will succeed—but quite good at forecasting that **some prior-generation models will fail** (as those models often do). In the logs, they quickly learn that certain Baselines (especially Gemma 2B and Haiku) are fragile on tasks like large primes or tricky modular exponentiation, and they confidently bet against them even while remaining more agnostic about the stronger models.

Because the number of truly hard items is small, these patterns are exploratory, but they suggest that LLM forecasters can quickly learn **model-specific "signatures" of weakness** and exploit them.

---

4. Discussion

4.1 What incentives did—and did not—do

The main takeaway of this pilot is not that fictional money makes models "smarter." The raw accuracy uplift is **positive but small** and not statistically significant at this sample size.

What incentives clearly *did* was make models **more legible**:

- In a standard classification setup, a "Yes" from the model looks identical whether its internal belief is 0.55 or 0.95.
- In this betting setup, the model compresses its internal probability distribution into a **single, actionable scalar**: the stake.