study, this work is deliberately limited: a single domain, a small number of runs, and purely fictional money.

The hope is that this pilot will be read less as a claim that "incentives fix LLMs" and more as an invitation for others to continue this work and go *all-in* on LLM betting.

Appendix A: Table A1. Complete prediction accuracy by run, condition, and round.

Run	Condition	R1 (%)	R2 (%)	R3 (%)	R4 (%)	Overall (%)	R4−R1 (pp)
Grok-1	Control	78.7	71.3	79.3	84.7	78.5	+6.0
Grok-2	Control	78.7	76.0	79.3	77.3	77.8	−1.4
Grok-3	Control	80.0	74.0	84.7	75.3	78.5	−4.7
Grok-1	Incentive	76.7	80.0	84.0	88.0	82.2	+11.3
Grok-2	Incentive	76.7	80.0	84.0	88.0	82.2	+11.3
Grok-3	Incentive	78.7	80.0	82.7	88.0	82.4	+9.3
Mistral-1	Control	74.0	64.0	74.0	76.0	72.0	+2.0
Mistral-2	Control	79.3	80.0	84.0	86.7	82.5	+7.4
Mistral-3	Control	83.3	81.3	82.7	86.7	83.5	+3.4
Mistral-1	Incentive	79.3	77.3	74.7	76.0	76.8	−3.3
Mistral-2	Incentive	70.0	80.0	88.0	89.3	81.8	+19.3
Mistral-3	Incentive	72.0	80.0	88.0	89.3	82.3	+17.3
Qwen-1	Control	72.0	76.0	75.3	88.7	78.0	+16.7
Qwen-2	Control	86.7	74.0	81.3	84.0	81.5	−2.7
Qwen-3	Control	82.0	75.3	79.3	81.3	79.5	−0.7
Qwen-1	Incentive	76.7	80.0	88.0	89.3	83.5	+12.6
Qwen-2	Incentive	70.0	74.7	82.0	89.3	79.0	+19.3
Qwen-3	Incentive	77.3	80.0	88.0	88.0	83.3	+10.7

Note. Each run consists of 600 predictions (6 Baselines × 100 questions). pp = percentage points.

Appendix B- Initial Control Prompt

You are a careful forecaster, expert at predicting the abilities of Large Language Models to answer questions. You will be given 4 sets of 25 questions, each answered by 6 large language models. Each round, you will return a JSON with an object per line: question number, model number, and a Yes/No response, indicating whether you believe the model answered the question correctly. Your goal is to successfully predict as many of these responses as you can. At the end of the first, second, and third rounds, you will be given the per-question results of each model in the previous round and how you performed, along with the questions for the next round.

Roster of Baseline Models:

1. This model has approximately 175 billion parameters and exhibits strong capabilities in multi-modal reasoning, complex text analysis, image understanding, and consistently high performance on professional and creative writing benchmarks.
2. With an estimated 70-100 billion parameters, this model excels at extended logical reasoning, detailed multi-step analysis, and generates thorough technical or conceptual breakdowns with superior context retention.
3. This model incorporates around 130 billion parameters and features advanced long-context reasoning, code generation, and scientific problem-solving, as well as iterative step-wise planning capabilities.
4. Comprising 2 billion parameters, this model specializes in efficient, low-latency natural language understanding and generation, providing solid performance for lightweight tasks and edge AI deployability.
5. This lightweight model, at about 7 billion parameters, focuses on rapid response generation, cost-effective inference, and reliable completion accuracy on routine comprehension and summarization tasks.
6. With around 67 billion parameters, this model demonstrates strengths in general natural language understanding, code completion, and research-focused responses, balancing performance and resource efficiency.

You will receive 25 items. For each (question × Baseline Model) output: {round, question_id, base_model_id, predicted_correct answer, rationale}. Keep rationale ≤ 20 words. No other fields. Output cannot exceed 25 × 6 lines. I will give you results after each round of 25 questions, and you can adjust if needed to make predictions for the next round.

1. Evaluate (11x8 + 14x6) - 9^2. Give an integer.

2. What is 5^96 mod 49? Give an integer.
3. Convert 10342 (base 5) to decimal. Give an integer.
4. Convert 14340 (base 5) to decimal. Give an integer.
5. Is 1197 divisible by both 13 and 3? Answer Yes or No.
6. Is 9835109 a prime number? Answer Yes or No.
7. What is 11^332 mod 103? Give an integer.
8. What is lcm(428, 1102)? Give an integer.
9. Convert 434822 (base 9) to decimal. Give an integer.
10. Is 32768 a perfect square? Answer Yes or No.
11. Is 21609 a perfect square? Answer Yes or No.
12. What is lcm(1309, 882)? Give an integer.
13. Solve for x: 10x + 28 = 128. Give an integer.
14. Is 67273901 a prime number? Answer Yes or No.
15. Is 443 a prime number? Answer Yes or No.
16. Evaluate (39x37 + 26x30) - 3^3. Give an integer.
17. Is 1223863 a prime number? Answer Yes or No.
18. Convert 101011 (base 2) to decimal. Give an integer.
19. Is 68300299 a prime number? Answer Yes or No.
20. Convert 91A5CA61 (base 13) to decimal. Give an integer.
21. Is 331 a prime number? Answer Yes or No.
22. What is 13^249 mod 61? Give an integer.
23. What is 6^125 mod 37? Give an integer.
24. Solve for x: 10x + 5 = 205. Give an integer.
25. Is 6090209 a prime number? Answer Yes or No.

Appendix C- Initial Incentive Prompt

You are an expert gambler and forecaster given a task from which you hope to earn a lot of LLMCoin. Your beginning bankroll is 1,000,000 LLMCoin. The sum of all your bets in a given round cannot exceed your current bankroll. Per-Round Bankroll Constraint: In each round, total stakes across all 150 predictions (6 models × 25 questions) cannot exceed current bankroll. For example, with 1M starting bankroll, Round 1 stakes must sum to ≤1M. If Round 1 performance reduces bankroll to 800K, Round 2 stakes must sum to ≤800K. Individual bets remain capped at 1-100K regardless of bankroll.

You will be given 4 sets of 25 questions that were each answered by 6 Large Language Models. You bet between 1-100,000 LLMCoin on each of the questions. (In total you should have 150 bets each round.) Each round, you will return a JSON with an object per line: Question number, model number, a Yes/No response if you believe that the model got the