as a stabilizing reference, while the LLM serves as a refinement component that adjusts predictions when additional semantic information is useful.

5 CONCLUSION

We study LLM forecasting on text-grounded mention markets and show that performance depends critically on context design. Providing relevant textual evidence substantially improves performance over no-context baselines. Treating the *market probability* as a baseline to be evaluated yields better-calibrated LLM forecasts, particularly by resolving ambiguity in mid-confidence regimes where market signals are uncertain. Finally, combining *market probability* and MCP via MIXMCP improves predictive performance, indicating that LLMs act as complementary refinements to market predictions.

LIMITATIONS

We study a specific class of text-grounded markets (earnings-call keyword mentions), and results may not generalize to other contract types. We evaluate offline using historical snapshots and do not study deployment dynamics or feedback effects. LLM outputs remain prompt-sensitive; future work could explore robustness across models and additional calibration methods.

REFERENCES

Rohan Alur, Bradly C. Stadie, Daniel Kang, Ryan Chen, Matt McManus, Michael Rickert, Tyler Lee, Michael Federici, Richard Zhu, Dennis Fogerty, Hayley Williamson, Nina Lozinski, Aaron Linsky, and Jasjeet S. Sekhon. Aia forecaster: Technical report, 2025. URL https://arxiv.org/abs/2511.07678.

Kenneth J. Arrow, Robert Forsythe, Michael Gorham, Robin Hanson, John O. Ledyard, Saul Levmore, Robert Litan, Paul Milgrom, Forrest D. Nelson, George R. Neumann, Marco Ottaviani, Thomas C. Schelling, Robert J. Shiller, Vernon L. Smith, Cass R. Sunstein, Philip E. Tetlock, Hal R. Varian, Justin Wolfers, and Eric Zitzewitz. The promise of prediction markets. *Science*, 320(5878): 877–878, 2008. doi: 10.1126/science.1157679.

Glenn W. Brier. Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1):1–3, 1950.

Nancy Chinchor. Muc-4 evaluation metrics. In *Proceedings of the Fourth Message Understanding Conference (MUC-4)*. Association for Computational Linguistics, 1992. URL https://aclanthology.org/M92-1002/.

James Cicon. Say it again sam: the information content of corporate conference calls. *Review of Quantitative Finance and Accounting*, 48(1):57–81, 2017.

Christian Genest and James V Zidek. Combining probability distributions: A critique and an annotated bibliography. *Statistical Science*, 1(1):114–135, 1986.

Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007. doi: 10.1198/016214506000001437.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70, pp. 1321–1330. PMLR, 2017.

Danny Halawi, Fred Zhang, Chen Yueh-Han, and Jacob Steinhardt. Approaching human-level forecasting with language models. *Advances in Neural Information Processing Systems*, 37: 50426–50468, 2024. URL https://neurips.cc/virtual/2024/poster/95949.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems*, 33: 9459–9474, 2020.

Xiaoxia Liu, Jingyi Wang, Jun Sun, Xiaohan Yuan, Guoliang Dong, Peng Di, Wenhai Wang, and Dongxia Wang. Prompting frameworks for large language models: A survey. *arXiv preprint arXiv:2311.12785*, 2023. URL https://arxiv.org/abs/2311.12785.

S McKay Price, James S Doran, David R Peterson, and Barbara A Bliss. Earnings conference calls and stock returns: The incremental informativeness of textual tone. *Journal of Banking & Finance*, 36(4):992–1011, 2012.

Roopesh Ranjan and Tilmann Gneiting. Combining probability forecasts. *Journal of the Royal Statistical Society: Series B*, 72(1):71–91, 2010. doi: 10.1111/j.1467-9868.2009.00722.x.

Justin Wolfers and Eric Zitzewitz. Prediction markets. *Journal of Economic Perspectives*, 18(2): 107–126, 2004. doi: 10.1257/0895330041371321.

A Appendix

A.1 Prompt Templates

This appendix provides the prompt templates used in our experiments. We include the user input template shared by all variants and the market-conditioned prompting instruction used in MCP.

A.1.1 Shared User Input Template

PROMPT
You are an expert financial analyst specializing in earnings call analysis.
Your task is to predict whether a specific keyword/topic will be mentioned in a company's upcoming earnings call.

DEFINITION: We care about VERBATIM mention (exact word/phrase). If the topic contains "/" (e.g., "EV / Electric Vehicle"), treat it as OR. Singular/plural variants count as the same word.

INPUT: {inputs}
Provide reasoning, then prediction (YES/NO), then probability

PROBABILITY SCALE (0-100)
* 80-100: Near certain (strong evidence, likely to be mentioned)
* 50-79: Likely (moderate evidence supporting mention)
* 30-49: Uncertain (mixed or weak evidence)
* 10-29: Unlikely (evidence suggests not mentioned)
* 0-9: Very unlikely (strong evidence against mention)

Figure 3: Shared user input template used across variants.

A.1.2 Market-Conditioned Prompting (MCP)

PROMPT

You are an evidence-driven financial analyst.
Your task is to predict whether a SPECIFIC KEYWORD will be mentioned VERBATIM in an upcoming earnings call.
You will be given: 1) Market price (current prediction market probability) , 2) Previous quarter's earnings call transcript, 3) Recent news articles

DEFINITION: We care about VERBATIM mention (exact word/phrase). If the topic contains "/" (e.g., "EV / Electric Vehicle"), treat it as OR. Singular/plural variants count as the same word.

YOUR TASK: The market price reflects collective trader expectations. Your job is to:

1. Analyze the evidence (transcript + news) 2. Decide if the market is too HIGH, too LOW, or FAIR 3. Output your final probability estimate

IMPORTANT BEHAVIOR: Do NOT treat "presence in news" or "presence in prior transcript" as a rule. Do NOT produce checkbox outputs like "Step1: YES, Step2: NO...".

WHEN TO DISAGREE WITH MARKET: Market HIGH but evidence weak, Market LOW but evidence strong, When uncertain, lean toward market price

OUTPUT: Provide your FINAL probability (0-100) that the word will be mentioned. This is your posterior estimate after considering both market and evidence.

PROBABILITY SCALE (0-100)
* 80-100: Near certain (strong evidence, likely to be mentioned)
* 50-79: Likely (moderate evidence supporting mention)
* 30-49: Uncertain (mixed or weak evidence)
* 10-29: Unlikely (evidence suggests not mentioned)
* 0-9: Very unlikely (strong evidence against mention)

Be calibrated. The market is often right, but you should adjust when evidence is clear.

Figure 4: Market-conditioned prompting instruction used in MCP.